

Q-1 Explain the fundamental differences between DDL, DML, and DQL commands in SQL. Provide one example for each type of command.

Ans– In SQL, commands are grouped based on **what kind of operation they perform on a database**. The three common categories are **DDL, DML, and DQL**.

1. DDL (Data Definition Language)

DDL commands define, modify, or remove the **structure of database objects** such as tables, schemas, or indexes.

Key characteristics:

- Work on database structure, not the actual data
- Changes are usually **auto-committed** (cannot be rolled back in most DBMS)

Example:

```
CREATE TABLE Students (
    StudentID INT,
    Name VARCHAR(50),
    Age INT
);
```

2. DML (Data Manipulation Language)

DML commands are used to **insert, update, delete, or modify data** stored in database tables.

Key characteristics:

- Work on the data inside tables
- Changes can usually be **rolled back** (transaction-controlled)

Example

```
INSERT INTO Students (StudentID, Name, Age)
```

```
VALUES (1, 'Suraj', 22);
```

3. DQL (Data Query Language)

Purpose:

DQL commands are used to **retrieve data** from one or more tables.

Key characteristics:

- Read-only operations
- Do not modify database structure or data

Example:

```
SELECT Name, Age  
FROM Students  
WHERE Age > 20;
```

Q-2 What is the purpose of SQL constraints? Name and describe three common types of constraints, providing a simple scenario where each would be useful.

Ans– SQL constraints are rules applied to table columns to **control the type, accuracy, and integrity of data** stored in a database. Their main purpose is to ensure **data consistency, correctness, and reliability** by preventing invalid data from being inserted or updated.

1. PRIMARY KEY Constraint

Description:

Ensures that each row in a table is **uniquely identifiable**. A primary key cannot contain **NULL values** and must be **unique**.

Useful scenario:

In a **Students** table, each student must have a unique student ID.

2. FOREIGN KEY Constraint

Description:

Ensures **referential integrity** by linking a column in one table to the primary key of another table. It prevents invalid references.

Useful scenario:

In an **Orders** table, every order must belong to an existing customer.

3. NOT NULL Constraint

Description:

Prevents a column from storing **NULL values**, ensuring that a value must always be provided.

Useful scenario:

A **User Registration** system where a username is mandatory.

Q-3 Explain the difference between LIMIT and OFFSET clauses in SQL. How would you use them together to retrieve the third page of results, assuming each page has 10 records?

ANS- Both **LIMIT** and **OFFSET** are used to control how many rows are returned and from where the result set starts, which is especially useful for pagination.

1. LIMIT

Purpose:

Specifies the **maximum number of rows** to return in the result set.

Example:

```
SELECT * FROM Employees
```

```
LIMIT 10;
```

Returns only the first 10 rows.

2. OFFSET

Purpose:

Specifies **how many rows to skip** before starting to return rows.

Example:

```
SELECT * FROM Employees
```

```
OFFSET 20;
```

Skips the first 20 rows and returns the remaining rows.

Q-4 What is a Common Table Expression (CTE) in SQL, and what are its main benefits? Provide a simple SQL example demonstrating its usage.

ANS– A Common Table Expression (CTE) is a **temporary, named result set** defined within the execution scope of a single SQL statement. It is created using the **WITH** keyword and can be referenced like a table in the main query.

Main Benefits of Using a CTE

1. **Improved Readability and Clarity**
Breaks complex queries into logical, easy-to-understand parts.
2. **Reusability Within a Query**
The same CTE can be referenced multiple times in the main query.
3. **Simplifies Complex Joins and Subqueries**
Replaces deeply nested subqueries with cleaner, more maintainable SQL.
4. **Supports Recursive Queries**
CTEs can be recursive, making them ideal for hierarchical data (e.g., organizational charts).

EXAMPLE:

```
WITH AvgSalary AS (
    SELECT AVG(Salary) AS avg_salary
    FROM Employees
)
SELECT Name, Salary
FROM Employees
WHERE Salary > (SELECT avg_salary FROM AvgSalary);
```

Q-5 Describe the concept of SQL Normalization and its primary goals. Briefly explain the first three normal forms (1NF, 2NF, 3NF).

ANS– SQL normalization is the process of organizing data in a relational database to reduce redundancy and improve data integrity. It involves dividing large tables into smaller, well-structured tables and defining relationships between them using keys.

Primary Goals of Normalization

1. **Eliminate data redundancy**
Avoid storing the same data in multiple places.
2. **Ensure data integrity and consistency**
Prevent update, insert, and delete anomalies.
3. **Improve database maintainability**
Make the structure easier to understand and modify.
4. **Efficient data storage**
Reduce unnecessary data duplication.

1. First Normal Form (1NF)

Rule:

- Each column must contain **atomic (indivisible) values**
- No repeating groups or multi-valued attributes
- Each record must be uniquely identifiable

Second Normal Form (2NF)

Rule:

- Must be in **1NF**
- All **non-key attributes** must be **fully dependent** on the **entire primary key**
- No **partial dependency** (applies to composite keys)

Third Normal Form (3NF)

Rule:

- Must be in **2NF**
- No **transitive dependency**
(Non-key attributes should not depend on other non-key attributes)

